

Art and Copyright after AI

Andrew Uroskie, 5/15/26

I. Three Shocks

Copyright has absorbed two structural shocks in the last three decades, and the second has not yet fully arrived. The first came in the 1990s, when digital reproduction collapsed distribution into copying and the marginal cost of a copy fell to zero. Legislators answered with the 1996 WIPO Copyright Treaty and its U.S. implementation as section 1201 of the Digital Millennium Copyright Act — statutes that turned the circumvention of any software lock into a federal crime, conscripting the criminal code into the enforcement of private business models (Giblin and Doctorow 2022, 25). The second shock has been underway since roughly 2020, and it arrives at a different order of magnitude. Where the 1990s problem concerned the unauthorized reproduction of individual works, the 2020s problem is the industrial ingestion of corpora numbering in the millions — corpora processed by neural networks whose outputs neither reproduce their training inputs nor remain orthogonal to them. Judge William Alsup’s June 2025 order in *Bartz v. Anthropic* describes the central library Anthropic assembled as containing over seven million pirated books retained “forever” as a general-purpose resource (Alsup 2025, 3). No prior legal regime had imagined a defendant of this kind.

The tempting move is to read the second shock as discontinuous with the first — as though generative AI presents a problem copyright was never designed to address. I want to argue the opposite. Generative AI does not break copyright so much as expose the ways in which it had always been misaligned: a Romantic-authorship paradigm awkwardly bolted onto industrial production; a fair-use doctrine that protects rights-holders’ markets more reliably than creators’ livelihoods; a visual-art jurisprudence doing the wrong work for two centuries because the art market itself runs on authenticity rather than copies; and an information economy whose dominant chokepoints are sustained less by copyright’s incentive logic than by the anti-circumvention statutes and platform contracts grafted onto it.

There is a third shock, older and quieter, that frames the others: the theoretical exhaustion of the Romantic-authorship paradigm itself, declared in French theory of the late 1960s and operationalized by appropriation art in the 1980s, but the law took no notice. The figure of the singular originating author remained the load-bearing fiction of copyright doctrine for another half-century, and remains it now. AI does not destroy copyright; it makes visible misalignments the digital turn had already produced, and that the avant-garde and structuralist critique had produced before that. The misalignment is the long story; AI is merely its latest visible chapter.

Theory Before Law: The Author Was Already Dead

Anglo-American copyright's "author" was forged in the eighteenth century and naturalized in the nineteenth. As Alexander Hartley puts it, the Romantic-authorship paradigm was itself "a 19th-century rearguard action" — a juridical fortification erected at precisely the moment when writers became market-dependent "literary workers," and when their claim to ownership of what they made required philosophical rather than economic grounding (Hartley 2024). The figure of the solitary genius producing original works from interior resources rationalized a system of contractually transferable property rights in literary expression. The doctrine made copyright possible but never described the actual conditions of literary or artistic production with any precision.

The late-1960s critique announced its intellectual exhaustion. Roland Barthes's "Death of the Author" (1968) and Michel Foucault's "What Is an Author?" (1969) furnished the textual-semiotic and the genealogical-juridical arguments for treating the author as a modern construction whose explanatory power was illusory. Barthes's closing formulation — "the birth of the reader must be at the cost of the death of the Author" (Barthes 1977, 148) — was less manifesto than observation: the text is "a tissue of quotations drawn from the innumerable centres of culture" (146), a multi-dimensional space whose unity sits with its destination rather than its origin. Foucault, complementarily, reread the author as a juridical-administrative category. Authors' names first appeared on texts not as a reward for creation but as a target for punishment. Ownership is liability's flipside — a point Hartley would later press into service for the AI moment.

By the late 1970s the avant-garde had operationalized the critique. The keystone text is Rosalind Krauss's "The Originality of the Avant-Garde" (1981), collected as the title essay of her 1985 volume of the same name. The modernist doctrine of originality — that the singular, unprecedented artwork grounds the artist-as-genius and underwrites the entire modernist art market — is itself a myth, and the grid is its emblem: modernist painting's founding originality gesture is in fact a repeatable structure whose logic is not development but repetition. "Once the neurosis is formed, we are specifically enjoined from thinking in terms of 'development,' and instead we speak of repetition" (Krauss 1985, 22). The volume's 1979 "Sculpture in the Expanded Field," collected in the same book, had already retroactively organized the 1970s as the historically necessary expansion of Minimalism into a four-term Klein-group field; its closing sentence names the binary the legal regime would later be forced to administer. "The postmodernist space of painting would obviously involve a similar expansion around a different set of terms from the pair architecture/landscape — a set that would probably turn on the opposition uniqueness/reproducibility" (Krauss 1985, 289). *Uniqueness/reproducibility* is the

unspoken matrix of every appropriation-art fair-use dispute the courts would spend the next four decades administering — from *Rogers v. Koons* through *Cariou v. Prince* and into *Warhol v. Goldsmith* — but the binary was named in print twenty years before the first of those rulings, in a book whose argument the law took no notice of. Douglas Crimp’s “Pictures” exhibition (Artists Space, 1977) supplied the curatorial apparatus: Sherrie Levine’s rephotography of Walker Evans, Cindy Sherman’s *Untitled Film Stills*, Robert Longo’s drawings, Barbara Kruger’s photo-text works, Louise Lawler’s documentation of art under collection. Levine’s rephotographs exposed Weston’s “author” as “dazzlingly multiple”; Sherman’s *Film Stills* were “a copy without an original” (Foster et al. 2016, 672–673). Krauss’s own contemporaneous 1985 reading of Levine — the canonical primary-source statement of appropriation as the postmodern de-origination of the photographic author — runs alongside (Krauss 1985, 168–170). By 1980 the avant-garde had its own machinery for displacing the Romantic-authorship paradigm; by 1985 the binary the law would spend forty years administering had already been named.

The 1985 volume’s most legally consequential passages show the doctrine of the original to be a juridical fiction the law has always quietly underwritten. Rodin’s *Gates of Hell*, juridically a “real original” by virtue of the Chambre des Députés’ 1978 authorization of twelve posthumous casts (Krauss 1985, 151), is in fact one of “multiple copies that exist *in the absence of an original*.” The issue of authenticity is equally problematic for each of the existing casts; it is only more conspicuously so for the most recent” (Krauss 1985, 152). Jean Chatelain, former director of the Louvre, supplies the doctrinal frame: “the special worth of an original edition does not come from an objective character of its originality ... since every edition is in itself an operation of reproducing a model. It arises from the agreements made by the edition’s author with the buyers” (Krauss 1985, 176). Originality, in the operative legal regime, is a contractual artifact. Even the matrix authorizing those casts was curatorially mediated: after Rodin’s 1917 stroke, the *Gates of Hell* plaster was reassembled not by him but by Léonce Bénédite of the Musée Rodin, who “took uncalled for and insensitive liberties” with arrangements “different than that which the artist wanted” (Krauss 1985, 188–189). What the 1978 casts faithfully reproduce is a curator’s posthumous edit. The Visual Artists Rights Act of 1990 would codify precisely that artifact — a “limited edition of 200 copies or fewer that are signed and consecutively numbered by the author” — as the statutory definition of visual art.

The law registered none of this. The figure displaced in theory by 1971 and in practice by 1980 remained the figure on which the legal regime depended: the 1976 Copyright Act, the 1990 Visual Artists Rights Act, the 1998 DMCA, the work-for-hire expansions, the international WIPO harmonization. The structural problem was nameable by 1971. It would take the next half-century before the law would be forced to look at it.

The Digital Turn: Code as Law

The first of the two technological shocks arrived in the 1990s. The legislative response migrated copyright's enforcement architecture from the work itself to the lock around it. The 1996 WIPO Treaty supplied the international framework; the 1998 DMCA, in particular its section 1201, the U.S. implementation. The conceptual frame was already in print. Lawrence Lessig's *Code: Version 2.0* (2006), revising the 1999 first edition, argued that "the software and hardware (i.e., the 'code' of cyberspace) that make cyberspace what it is also regulate cyberspace as it is" (Lessig 2006, 5). The slogan — "code is law" — has a precise content: architecture exerts the same regulatory force as legal statute, and the state, regulating directly through commands, also regulates indirectly by adjusting any of the others.

Section 1201 made it a federal crime, punishable by five years' imprisonment and a \$500,000 fine for a first offense, to bypass any software lock on a copyrighted work, regardless of whether the underlying use was lawful (Giblin and Doctorow 2022, 25). Rebecca Giblin and Cory Doctorow's reading of this clause is the indispensable one. Section 1201, they argue, created "felony contempt of business model": a new cause of action available to anyone who could wrap a copyrighted work in DRM, regardless of whether the use the lock prevented was infringing in the first place (Giblin and Doctorow 2022, 38). The cause of action does not require copyright infringement; anything wrapped in DRM generates a new federal crime when the wrapper is bypassed. Amazon's 89-percent share of the ebook market by 2018 was secured not by superior copyright protection but by the legal sanctity of the Kindle DRM publishers had themselves insisted on (Giblin and Doctorow 2022, 32). Apple's 30-percent App Store toll, enforced by the same statute, demonstrated that the lock was a tollgate before it was a security measure (Giblin and Doctorow 2022, 121). Copyright's exclusive right to control reproduction — the incentive engine that justified the entire system — had been operationally subordinated to DRM-secured chokepoints whose owners were neither the authors copyright purported to protect nor the public it was supposed eventually to serve.

The cultural-political counter-movement is best traced through Lessig's *Free Culture* (2004) and *Remix* (2008), and through Aaron Swartz's activist counter-voice — the 2008 Guerilla Open Access Manifesto, the JSTOR action, the federal indictment that would have consumed Swartz's fortune and freedom (Swartz 2016). The cultural critique was won; permission culture as a name for the regime *Free Culture* described was accepted across the left and much of the libertarian right. The legal infrastructure was not changed. The chokepoints the section 1201 regime had built were absorbed by the platforms that grew up on top of them. DRM was displaced upward into terms-of-service over which users had no negotiation. Streaming-revenue distribution shows the operational outcome unchanged: by 2002, on the *Chokepoint Capitalism* figures, just seven

percent of recording-industry revenues reached performers and songwriters (Giblin and Doctorow 2022, 54). Remix as a content format was absorbed; the underlying enforcement architecture was not contested. AI training corpora, when they arrived, would ignore even the terms-of-service layer entirely.

Luciano Floridi's 1996 "Brave.Net.World," written just before the DMCA, identified the converging informational pathology. Disinformation, Floridi argued, is endogenous to any information-management system, "from the manuscript tradition to the card index of a library" (Floridi 1996, 4). Three trends would intensify it on the digital substrate: the rise of "provusers" who blur the provider/user distinction; the integration of media types that had previously inhibited unilateral capture; and the "self-disinformation" enabled by interactive filtering. Read with thirty years' hindsight, Floridi's "provuser" describes precisely the figure AI training corpora would come to feed on. The DMCA had already laundered authorship into platform control. AI would launder it into statistical weight.

The Doctrinal Architecture

The courts were not idle through any of this, and the doctrinal architecture they built deserves to be read as architecture rather than as a sequence of political outcomes. The fair-use inquiry is statutory — four factors enumerated in § 107 of the Copyright Act — but the operative doctrine has been built case by case, and the cases that matter most are the ones that defined the structural questions later cases would answer. Three such decisions stand at the load-bearing points of the structure, and the arc among them is coherent.

Sony Corp. of America v. Universal City Studios (1984) established that a copying technology "capable of substantial noninfringing uses" cannot be banned as contributorily infringing. The Betamax decision is the structural answer to a recurring question: when a technology enables both lawful and unlawful uses, the law does not foreclose the technology to prevent the unlawful ones. The principle would shape every subsequent technology-platform dispute, from peer-to-peer file sharing to AI training corpora.

Campbell v. Acuff-Rose Music (1994) is the hinge from which transformative-use doctrine descends. Justice Souter, adopting Pierre Leval's formulation, held that the central inquiry under factor one is whether the new work "merely 'supersede[s] the objects' of the original creation ... or instead adds something new, with a further purpose or different character, altering the first with new expression, meaning, or message" (*Campbell* 1994, 579). The opinion repudiated the earlier dictum that "every commercial use ... is presumptively ... unfair" and opened conceptual space for fair use in cases where a secondary work transformed rather than substituted for its

source. *Campbell* is the doctrinal seed of every subsequent appropriation-art fair-use case, and of the “transformative” defense AI training would later invoke.

Andy Warhol Foundation v. Goldsmith (2023), three decades after *Campbell*, contracted the doctrine *Campbell* had expanded. The Supreme Court’s first fair-use ruling on visual art concerned the Warhol Foundation’s 2016 commercial license to *Condé Nast* — for \$10,000 — of “Orange Prince,” one work in Warhol’s 1984 sixteen-work Prince Series derived from Lynn Goldsmith’s 1981 studio portrait. Justice Sotomayor, writing for a 7–2 majority, introduced what is best called the **specific-use test**: the relevant inquiry under factor one is the purpose of the *challenged use* — here, the Foundation’s commercial license to a magazine — not the abstract aesthetic transformation of the underlying photograph. “The same copying may be fair when used for one purpose but not another” (*Warhol Foundation* 2023, 20). Where an original work and a secondary use share the same or highly similar purposes, and the secondary use is commercial, the first factor is likely to weigh against fair use absent some other justification. Justice Kagan, joined by Chief Justice Roberts, dissented at length, defending “appropriation, mimicry, quotation, allusion and sublimated collaboration” as “a sine qua non of the creative act” (*Warhol Foundation* 2023, 71) and warning that the majority’s test would “stifle creativity of every sort” (*Warhol Foundation* 2023, 81).

What the opinion declines to address — what no copyright opinion could address — is the third party at the dispute’s center. Jennifer Rothman maps how all three courts erased Prince Rogers Nelson procedurally: the district court dissected “contours of Prince’s face,” “his neckline,” “Prince’s torso”; the Second Circuit framed the conflict as one between “artists each seeking to profit from his or her own creative efforts” with Prince entirely unrecognized; the Supreme Court “treats him like a dress-up doll, cutting out Prince’s head from the Warhol print and pasting it on to the original Goldsmith photograph” (Rothman 2024, 53). The licensing logic the Court invoked to vindicate Goldsmith, Rothman argues, could equally have been marshaled in Prince’s favor under the right-of-publicity branch she and Robert Post term the “right of commercial value” (Rothman 2024, 53). Copyright’s structural blindness to the subject is the doctrine’s load-bearing absence: a regime that adjudicates between photographer and appropriator while erasing the person whose image gives both works their commercial value.

The arc — *Sony*’s safe harbor for technology, *Campbell*’s transformative-use opening, *Warhol*’s specific-use contraction — is not a pendulum. It is a coherent institutional response to a recurring problem: how to administer an exclusive right whose scope is constantly being tested by uses the original statute did not imagine. Each decision answered a structural question the prior ones had left open. *Warhol* did not undo *Campbell*; it specified what *Campbell* had left ambiguous — that aesthetic transformation is necessary but not sufficient, and that fair use must be administered against the particular market the rights-holder credibly occupies. What *Warhol* could not

anticipate was the case that would arrive eighteen months later, in which the “challenged use” was the training of a foundation model on seven million books.

Art Was Always Outside

Visual art has been quietly demonstrating copyright’s misfit with creative production for two centuries. Noam Elcott and Amy Adler, in their editors’ introduction to the *Grey Room* special issue on “Art beyond Copyright,” put the empirical observation plainly: “Visual artists today make little if any of their money from copies” (Elcott and Adler 2024, 11). Even the Andy Warhol Foundation’s licensing income “pales in comparison to the multi-million-dollar prices fetched by individual Warhol canvases” (Elcott and Adler 2024, 12). Copyright’s exclusive right to make copies — the right that justifies the entire system on incentive grounds — “does not provide an economic incentive to visual artists” (Elcott and Adler 2024, 13). What provides the incentive is a parallel norm Elcott and Adler call “authenticity”: a copy of a Warhol, however faithfully executed, “has effectively no market value — unless executed by someone like Elaine Sturtevant or Deborah Kass” (Elcott and Adler 2024, 13). Artist signatures function as trademarks; artworks function as luxury goods.

The Visual Artists Rights Act of 1990 codifies the misalignment without quite acknowledging it. VARA defines “visual art” as “a painting, drawing, print, or sculpture, existing in a single copy, or in a limited edition of 200 copies or fewer that are signed and consecutively numbered by the author” (Elcott and Adler 2024, 11, citing 17 U.S.C. § 106A). The definition reads as “perverse and startlingly retrograde in relation to generations of avant-garde attacks on authorship, aura, originality.” Yet, Elcott and Adler argue, it discloses “an unwitting truth”: the artist, in the operative legal definition, is precisely “an author who makes little to no money from mass copies.” The legal system has been forced to acknowledge what its own incentive theory cannot accommodate.

The cost of the misalignment falls hardest on those least equipped to absorb it. A 2014 College Art Association study found art historians paying licensing fees from their own pockets — in one case, \$20,000 for a single book — while a majority of arts editors had “avoided or abandoned a project for copyright reasons” (Elcott and Adler 2024, 15). Visual-art litigation, the editors note, “appears nowhere in the case law for the simple reason that copyright wielded in this realm never rises to the level of litigation”: the cases that do reach federal courts feature “the top 0.001 percent of artists” (Elcott and Adler 2024, 15). *Warhol v. Goldsmith* is symptomatic — the first Supreme Court fair-use case on visual art is a one-in-a-million dispute that bears little relation to the conditions of the field whose practice it now constrains.

Stefano Basilico's 2004 *Cut: Film as Found Object* catalogue documents the longer history the AI debate has tended to obscure. Found-footage practice — Esfir Shub's *Fall of the Romanov Dynasty* (1927), Joseph Cornell's *Rose Hobart*, Bruce Conner's *A Movie*, Christian Marclay's *Telephones* — operated outside the legal permissions regime for the better part of a century. "Art has never been a permission culture, even though strategies of quotation, appropriation, pastiche, and outright copying are virtually as old as the practice of art making itself" (Basilico 2004, 7). Lessig's contribution to the same catalogue diagnoses the contemporary regime as a "permission culture" against which the avant-garde's practice constitutes a sustained refusal (Basilico 2004, 47). The avant-garde tradition the AI debate now invokes by analogy is one the law has been quietly tolerating for decades, on the implicit understanding that to litigate it at scale would be to confess copyright's incompatibility with how art actually works.

The misalignment runs deeper than the editors' introduction states. Amy Adler's individual contribution to the same *Grey Room* issue diagnoses moral rights, codified by VARA, as what Peter Baldwin called "the anti-copyright" — a personality-theory tradition grounded in continental civil law and fundamentally hostile to the utilitarian-economic philosophy U.S. copyright defends (Adler 2024, 63). *Mazer v. Stein* (1954) articulated the latter as "an economic philosophy ... that encouragement of individual effort by personal gain is the best way to advance public welfare" (Adler 2024, 64); VARA's noneconomic, dignitary protections collide with the regime into which Congress slotted them. Because VARA's scope is so narrow, artists increasingly deploy copyright itself "as an off-label mechanism ... to pursue moral rights types of concerns" (Adler 2024, 65) — using doctrine designed for economic incentives to vindicate dignitary harms it was never built to recognize. Xiyin Tang, writing alongside, identifies a parallel layer: the Trademark Counterfeiting Act of 1984 introduced criminal liability explicitly modeled on art forgery, on the rationale that "trademarks play a similar role in commerce as signatures" (Tang 2024, 75). Tang reads Krauss on Rodin — "recognized as authentic only those bronze casts he had authorized" (Tang 2024, 75; cf. Krauss 1985, 152) — as already operating the brand-signature logic trademark law would later import. The apparatus governing visual art is distributed across at least four regimes — copyright, VARA, publicity rights, trademark — each with a different theory of what an artwork is. None describes the field with much precision.

The Generative Turn: *Bartz* and Hartley's Twofold Problem

The 2020s shock arrived at the level of doctrine on June 23, 2025, when Judge William Alsup issued the summary-judgment order in *Bartz v. Anthropic PBC*. The most economical statement of where the input side of copyright now stands is Alexander Hartley's "To Whom Does the World Belong?", published in Evgeny Morozov's *AI Futures* the previous year. Hartley splits the

AI copyright problem in two: an *input* problem — the thirteen-trillion-token training corpus of mostly copyrighted text on which any frontier model depends — and an *output* problem, which Hartley schematizes as the six candidate owners of any AI-generated work: developer, licensee, training-data holders, prompter, the public domain, or the AI itself.

Bartz settles part of the input question and leaves the output question almost wholly untouched. Alsup’s order distinguishes three uses of the same copyrighted books: training copies used to develop the LLM weights; print-to-digital format conversions of books Anthropic had purchased and physically scanned; and pirated copies retained, beyond the act of training itself, as a general-purpose library accessible to Anthropic’s engineers (Alsup 2025, 9–10). On the first, the court holds the training use “quintessentially transformative,” analogous to “a reader aspiring to be a writer” who “trained upon works not to race ahead and replicate or supplant them” (Alsup 2025, 13). On the second, the court reaches back to *Sony*: format conversion is a space-shift, “one copy replacing another, not multiplying copies or sharing them outside the company” (Alsup 2025, 16). On the third — the seven-million-volume central library Anthropic assembled from Books3, Library Genesis, and Pirate Library Mirror — summary judgment is denied. “Building a central library of works to be available for any number of further uses was itself the use for which Anthropic acquired these copies”; the library is not subsumed under the transformative purpose of training and substitutes for paid copies in a general-purpose research collection (Alsup 2025, 19).

The facts the court accepts as undisputed are themselves the story. Anthropic co-founder Ben Mann downloaded Books3 (196,640 volumes), then at least five million books from Library Genesis, then at least two million from Pirate Library Mirror — over seven million pirated books retained “forever” as a general-purpose resource accessible to hundreds of engineers (Alsup 2025, 3–4). Dario Amodei is on record preferring piracy to licensing “to avoid legal/practice/business slog” (Alsup 2025, 2–3). The training is fair; the library is not. The chokepoint architecture section III traced — the DMCA’s lock-protection layer, contractually extended through platform terms of service — Anthropic simply bypassed at scale.

The doctrinal stakes extend beyond *Bartz*. Kate Crawford and Jason Schultz, writing in *Grey Room* the same year, argue that generative AI induces what they call “an epistemic crisis of law,” destabilizing three foundational concepts simultaneously: authorship, expressiveness, and agency (Crawford and Schultz 2024, 56). On *authorship*, the U.S. Copyright Office’s categorization of generative outputs as the work of “machine authors” — ineligible for copyright on the *Burrow-Giles / Naruto* line — has produced “the largest experiment in ‘art beyond copyright’ to date,” billions of works in the public domain by operation of doctrine (Crawford and Schultz 2024, 57). On *expressiveness*, the idea-expression distinction of *Baker v. Selden* (1879) has been chronically unstable — Learned Hand’s 1930 admission in *Nichols v. Universal Pictures* that “nobody has

ever been able to fix that boundary, and nobody ever can.” Generative systems radicalize the instability: prompts are ideas, outputs are expressions, and a single prompt yields millions of possible expressions. The interface is “essentially *an idea-expression translation machine* ... an epistemic skein that copyright law may be hopelessly unable to untangle, especially at scale” (Crawford and Schultz 2024, 59); the echo of Sol LeWitt’s “the idea becomes a machine that makes the art” is deliberate. On *agency*, the volitional-conduct doctrine of *Cartoon Network v. Cablevision* (2d Cir. 2008) located infringement-causation in the human actor who pressed the button; distributed neural-net production leaves no such actor identifiable — “not the researchers who assembled the training data, not the company that trained the model, not the user who wrote the prompt” specifically directed or could have predicted any particular infringing output (Crawford and Schultz 2024, 60). Three load-bearing doctrines fail simultaneously.

The output side is where Hartley’s historical recovery is sharpest. *Urantia Foundation v. Maaherra* (9th Cir. 1997) concerned William S. Sadler’s 1955 *Urantia Book*, transcribed from claimed alien communications and held to qualify for copyright on the strength of the Forum’s “choice of questions” — a level of human creativity Hartley quotes the court characterizing as “extremely low” but sufficient. “Many people would say Sadler and his friends were delusional,” Hartley writes; “today we might call them prompt engineers.” The case has been cited in *Thaler v. Perlmutter*, which placed Stephen Thaler’s Creativity Machine output “A Recent Entrance to Paradise” in the public domain on the ground that no human author existed. *Moffatt v. Air Canada*, in which a chatbot hallucinated a bereavement policy and the airline was held liable for it, demonstrates the Foucauldian principle Hartley presses: authors’ names first appeared on books not as a reward for creation but as a target for punishment. Ownership is liability’s flipside; the six candidate owners are six candidate defendants.

Hartley’s prescription is structural rather than doctrinal. Shorten copyright terms for AI-assisted works; assign autonomously generated elements to the public domain; consider whether private patents on AI systems themselves should be permitted at all, on Charles Jennings’s analogy to Truman’s 1946 Atomic Energy Commission, under which the federal government became *owner* but not operator of atomic technology. The stakes are not modest: ninety percent of recent U.S. economic growth and sixty-five percent of S&P 500 valuation is now intellectual property; AI services may reach a trillion dollars in annual revenue by 2032 (Hartley 2024). What *Bartz* brackets but cannot dissolve is the technical admission its own factual record contains: the trained model retains “compressed” copies of its training works, having “memorized” them almost verbatim, with the verbatim regurgitation filtered before reaching end users (Alsup 2025, 7). The corpus is *inside* the model, even when the law treats outputs as orthogonal to inputs. Barthes’s 1968 essay may at last be vindicated — not on literary-theoretical grounds, but on the

brute juridical ground that the figure of the author no longer corresponds to anything the law can find.

Three Empires and the Limits of the Visible

The *Bartz* ruling is not a global precedent; it is a national settlement inside a larger contest. Anu Bradford's *Digital Empires* organizes that contest around three competing regulatory models — American market-driven, Chinese state-driven, European rights-driven — each functioning as “an informal empire of the internet era” projecting “economic, military, and cultural power across [its] borders” (Bradford 2023, 13). The United States exercises *private* power through tech-company reach; China exercises *infrastructure* power through material exports; the EU exercises *regulatory* power through what Bradford has elsewhere theorized as the Brussels Effect (Bradford 2023, 29). When forced to balance competing imperatives, each falls back: the U.S. on pro-market restraint, China on the government's interest, the EU on the rights of digital citizens (Bradford 2023, 18). The EU AI Act and the Copyright Directive's text-and-data-mining provisions require both opt-out and training-data transparency; China treats training corpora as state-strategic within the “Made in China 2025” framework (Bradford 2023, 87–91). The Alsup settlement governs only one jurisdiction in this three-cornered field.

Tim Berners-Lee's *This Is for Everyone* (2025) offers a protocol-level alternative cutting across all three empires. The core proposal on AI training is opt-in rather than opt-out, with compensation modeled on YouTube's Content ID system; cryptographic chain-of-custody metadata via the C2PA standard; and a data-sovereignty substrate built on the Solid protocol and Personal Online Data Stores (Berners-Lee 2025, ch. 16). “Imagine if the same incentives for engagement that have made social media so toxic are allowed to dictate the development of AI” (Berners-Lee 2025, ch. 17). Brian J. Chen's observation in the same *Boston Review* forum sharpens the geopolitical point: TSMC, the foundry on which the entire AI compute stack depends, “from day one ... wasn't really a private business: it was a project of the Taiwanese state” (Chen 2024). The combined market capitalization of the five largest U.S. tech companies — over seven trillion dollars — was built on training data scraped from a global commons whose contributors had no meaningful capacity to opt out (Bradford 2023, 8). What U.S. doctrine calls fair use, the geopolitics calls sovereign expropriation by another name.

Noam Elcott's essay on Arthur Jafa's *BG* (2024) performs the diagnostic operation *Bartz* cannot. *BG* digitally replaces the white pimp, bouncer, and mafioso in the climactic shoot-up of Scorsese's *Taxi Driver* with Black actors, restoring the racial animus present in Paul Schrader's original screenplay and suppressed by the finished film. Jafa “cleared no rights for *BG*” — Elcott

calls this “a striking move in 2024” in the wake of *Warhol Foundation v. Goldsmith* (Elcott 2025). Schrader publicly praised the result. The work depends on the law’s selective blindness — the same blindness Judith Butler diagnosed in the Rodney King trial: “to the extent that there is a racist organization and disposition of the visible, it will work to circumscribe what qualifies as visual evidence” (Butler, quoted in Elcott 2025). The same appropriation that earns Jafa praise would, in another body or another genre, generate litigation. The permission regime is selective in ways the doctrine cannot articulate.

Florida’s endogenous-disinformation diagnosis has a structural correlate in the legal information system. Fair-use balancing is, by design, fact-intensive; case law reaches only the top 0.001 percent of artistic practice; the doctrine has been built almost entirely from cases brought by Warhol, Koons, and Prince. Generative AI compounds these conditions with a fourth: endogenous attribution-loss. In a corpus of seven million books, authorial contribution disappears into statistical weights. The corpus is inside the model. The law’s bifurcation of inputs from outputs is a convenience the technology will not sustain indefinitely.

After Romantic Authorship Today

Three propositions:

First: the Romantic-authorship paradigm is no longer load-bearing. Barthes’s death of the author, Krauss’s death of the original, and Hartley’s death of the prompter-as-author converge on a single conclusion the law has spent sixty years deferring. The figure survived the rise of corporate copyright, the work-for-hire doctrine, and the industrial aggregation of publishing — because in each of those cases the author could still be *located*, even if displaced into a contractual signature on a record label’s standard form. The training of a foundation model on thirteen trillion tokens admits no such localization. The author is no longer displaced but dissolved. What replaces the figure — a system-level liability rule on the *Moffatt v. Air Canada* model, a public-domain default for autonomously generated outputs, a compulsory-licensing regime on the Content ID model Berners-Lee proposes — is the question. What the law cannot do is restore the figure to its prior office.

Second: copyright was already the wrong instrument for protecting creators’ livelihoods, and AI has made the wrongness operational rather than theoretical. The DMCA’s section 1201 turned copyright into a chokepoint-enforcement statute. The VARA misfit demonstrates that visual art runs on authenticity rather than copies. The share of streaming revenue reaching songwriters sits in the single digits not because copyright is too weak but because copyright was never the binding constraint (Giblin and Doctorow 2022, 54). The remedies on the table — automatic

reversion of rights to creators, residual remuneration enforceable downstream, statutory floors for creative labor, mandated interoperability — operate *alongside* copyright, not through it. AI sharpens the case for these remedies precisely by demonstrating that whatever copyright was supposed to do for individual creators, it cannot do for creators whose work has been ingested into a corpus too large for the doctrine to track.

Third: the geopolitics of training data has overtaken the politics of copyright. The next regulatory settlements will be made at the level of regime competition — American private power, Chinese infrastructure power, European regulatory power — rather than through doctrinal refinement inside any one jurisdiction. Berners-Lee’s protocol-level proposal offers a fourth path that does not depend on the three empires reconciling: opt-in training data, cryptographic provenance, user-controlled data stores. Whether any of the four paths prevails is undecided.

The Romantic author has been declared dead by every era since 1968 — by Barthes and Foucault in theory, by the Pictures Generation in practice, by the DMCA in enforcement, and now by *Bartz* in doctrinal coherence. What replaces that figure, and on whose terms, is the question every era has handed forward unanswered. AI has raised the stakes such that the answer can no longer be deferred.

* * *

Bibliography

Adler, Amy. “Moral Rights: The Anti-copyright.” *Grey Room* 94 (Winter 2024): 63–66.

Barthes, Roland. *Image Music Text*. Translated by Stephen Heath. New York: Hill and Wang, 1977.

Basilico, Stefano, ed. *Cut: Film as Found Object in Contemporary Video*. Milwaukee: Milwaukee Art Museum, 2004.

Berners-Lee, Tim, with Stephen Witt. *This Is for Everyone: The Unfinished Story of the World Wide Web*. New York: Farrar, Straus and Giroux, 2025.

Bradford, Anu. *Digital Empires: The Global Battle to Regulate Technology*. New York: Oxford University Press, 2023.

Chen, Brian J. “Semiconductor Island.” In *AI Futures*, edited by Evgeny Morozov. Boston Review Forum 32. Cambridge, MA: Boston Review, 2024.

Crawford, Kate, and Jason Schultz. “The Work of Copyright Law in the Age of Generative AI.” *Grey Room* 94 (Winter 2024): 56–62.

Elcott, Noam M. “Arthur Jafa’s *BG*, or, Seeing Race and Visual Truth in the Age of AI.” *Grey Room* 100 (Summer 2025): 62–89.

Elcott, Noam M., and Amy Adler. “Editors’ Introduction: Art beyond Copyright.” *Grey Room* 94 (Winter 2024): 6–21.

Floridi, Luciano. “Brave.Net.World: The Internet as a Disinformation Superhighway?” *The Electronic Library* 14, no. 6 (1996): 509–514.

Foster, Hal, Rosalind Krauss, Yve-Alain Bois, Benjamin H. D. Buchloh, and David Joselit. *Art Since 1900: Modernism, Antimodernism, Postmodernism*. 3rd ed. London: Thames & Hudson, 2016.

Giblin, Rebecca, and Cory Doctorow. *Chokepoint Capitalism: How Big Tech and Big Content Captured Creative Labor Markets and How We’ll Win Them Back*. Boston: Beacon Press, 2022.

Hartley, Alexander. “To Whom Does the World Belong?” In *AI Futures*, edited by Evgeny Morozov. Boston Review Forum 32. Cambridge, MA: Boston Review, 2024.

Krauss, Rosalind E. *The Originality of the Avant-Garde and Other Modernist Myths*. Cambridge, MA: MIT Press, 1985. Collects “Sculpture in the Expanded Field” (originally *October* 8, Spring 1979, 30–44), “The Originality of the Avant-Garde” (originally *October* 18, Fall 1981, 47–66), and “Sincerely Yours.”

Lessig, Lawrence. *Code: Version 2.0*. New York: Basic Books, 2006.

———. *Free Culture: How Big Media Uses Technology and the Law to Lock Down Culture and Control Creativity*. New York: Penguin Press, 2004.

———. *Remix: Making Art and Commerce Thrive in the Hybrid Economy*. New York: Penguin Press, 2008.

Morozov, Evgeny, ed. *AI Futures*. Boston Review Forum 32. Cambridge, MA: Boston Review, 2024.

Rothman, Jennifer E. “The Absent Prince: Reflections on Personality Rights and *Andy Warhol Foundation v. Goldsmith*.” *Grey Room* 94 (Winter 2024): 51–55.

Swartz, Aaron. *The Boy Who Could Change the World: The Writings of Aaron Swartz*. New York: The New Press, 2016.

Tang, Xiyin. “Trademark Law and the Contingent Art Object.” *Grey Room* 94 (Winter 2024): 73–76.

Caselaw

Andy Warhol Foundation for the Visual Arts, Inc. v. Goldsmith, 598 U.S. 508 (2023).

Baker v. Selden, 101 U.S. 99 (1879).

Bartz v. Anthropic PBC, No. 3:24-cv-05417-WHA (N.D. Cal. June 23, 2025) (order on cross-motions for summary judgment).

Burrow-Giles Lithographic Co. v. Sarony, 111 U.S. 53 (1884).

Campbell v. Acuff-Rose Music, Inc., 510 U.S. 569 (1994).

Cartoon Network LP v. CSC Holdings, Inc., 536 F.3d 121 (2d Cir. 2008).

Mazer v. Stein, 347 U.S. 201 (1954).

Moffatt v. Air Canada, 2024 BCCRT 149 (Civ. Resol. Trib. Feb. 14, 2024).

Naruto v. Slater, 888 F.3d 418 (9th Cir. 2018).

Nichols v. Universal Pictures Corp., 45 F.2d 119 (2d Cir. 1930).

Sony Corp. of America v. Universal City Studios, Inc., 464 U.S. 417 (1984).

Thaler v. Perlmutter, 687 F. Supp. 3d 140 (D.D.C. 2023).

Urantia Foundation v. Maaherra, 114 F.3d 955 (9th Cir. 1997).